

CS7641 Assignment 3

Unsupervised Learning and Dimensionality Reduction

1 Assignment Weight

The assignment is worth 20% of the total points.

Read everything below carefully as this assignment has changed term-over-term.

2 Objective

Now it is time to explore unsupervised learning algorithms. This part of the assignment asks you to use some of the clustering and dimensionality reduction algorithms we've looked at in class and to revisit earlier assignments. The goal is for you to think about how these algorithms are the same as, different from, and interact with your earlier work.

The same ground rules apply for programming languages and libraries. You may program in any language that you wish insofar as you feel the need to program. As always, *it is your responsibility* to make sure that we can actually recreate your narrative, if necessary.

Please note, this class implements changes to the assignments term-over-term as we are calibrating the course incrementally. Please read through everything, even if you are submitting work from a previous semester as the requirements will likely have changed.

3 Procedure

3.1 The Problems Given to You

You are to implement six algorithms.

The first two are clustering algorithms. You can choose your own measures of distance/similarity. Justify your choices.

- Expectation Maximization
- Clustering Algorithm of Your Choice

The last four are dimensionality reduction algorithms:

Linear:

- PCA
- ICA
- Randomized Projections

Non-linear:

- Manifold Learning Algorithm of Your Choice

You are to run several experiments with the goal of disseminating how dimensionality reduction affects your data. Come up with at least two datasets. You can use the datasets from previous assignments. Develop hypotheses based on your datasets and the following exploration. These hypotheses should be well-posed and grounded in theory from the lectures and readings.

3.2 Exploration

The following should comprise your exploration.

1. Apply the clustering algorithms on the datasets. You will report on each of the clustering algorithms for each dataset, resulting in 4 demonstrations.
2. Apply the dimensionality reduction algorithms on the datasets. You will report on each of the dimensionality reduction algorithms, resulting in 8 demonstrations.
3. Re-apply the clustering algorithms on the set of dimensionality reduction datasets. This will result in 16 combinations of results of datasets, dimensionality reduction, and clustering methods. You should look at the full scope of the results and note how they might pertain to your hypotheses. In particular, focus on more interesting findings. You will be reporting one of your clustering algorithms on your datasets with one linear method (PCA, ICA, or RP) and the manifold learning algorithm of your choice, resulting in 4 total demonstrations between the two datasets. Justification will be especially important as space is limited in the report.
4. Choose one of your datasets. Re-run your neural network learner from Assignment #1 with each of the dimensionality reduction algorithms applied. You will report on a different linear method from Step 3 (PCA, ICA, or RP) and on the manifold learning algorithm of your choice, resulting in 2 total demonstrations. Justification will be especially important as space is limited in the report.
5. Using the same dataset as Step 4, use both previously generated clusters from Step 1 as new features in your dataset. Again, rerun your neural network learner on the newly projected data and note the findings. You will report on each of the clustering algorithms, resulting in 2 demonstrations. Justification will be especially important as space is limited in the report.

Analysis writeup is limited to 8 pages. This is different from the first set of assignments. The page limit does include your citations. Anything past 8 pages will not be read. As a final check during your submission process, download the submission to double check everything looks correct on Canvas.

Please keep your analysis as concise while still covering the requirements of the assignment.

3.3 Acceptable Libraries

Here are a few **examples** of acceptable libraries. You can use other libraries as long as they fulfill the conditions mentioned above.

Machine learning libraries:

- scikit-learn (python)
- Weka (java)
- e1071 (R)
- ML toolbox (matlab)

Plotting:

- matplotlib (python)
- seaborn (python)
- yellowbrick (python)
- ggplot2 (R)

4 Submission Details

You must submit:

- A file named README.txt containing instructions for running your code (see note below)
- A file named yourgtaccount-analysis.pdf containing your writeup (GT account is what you log in with, not your all-digits ID)

The file yourgtaccount-analysis.pdf should contain:

- Brief description of your datasets, and hypotheses you want to highlight in your report.
- Explanations of methods. This is your opportunity to demonstrate nuances needed to support your hypotheses.
- Grounded descriptions of resulting clusters. Support descriptions with data-driven evidence.
- Analyses of your results. Why did you get the clusters you did? Do they make "sense"? If you used data that already had labels (for example data from a classification problem from assignment #1) did the clusters line up with the labels? Do they otherwise line up naturally? Why or why not? Compare and contrast the different algorithms. What sort of changes might you make to each of those algorithms to improve performance? How much performance was due to the problems you chose? Be creative and think of as many questions you can, and as many answers as you can. Take care to justify your analysis with data explicitly.
- Can you describe how the data looks in the new spaces you created with the various dimensionality reduction algorithms? For PCA, what is the distribution of eigenvalues? For ICA, how kurtotic are the distributions? Do the projection axes for ICA seem to capture anything "meaningful"? Assuming you only generate k projections (i.e., performing dimensionality reduction), how well is the data reconstructed by the randomized projections? How much variation did you get when you re-ran your random projections several times? How does noise affect each algorithm? What is the rank of your data? Can you describe how colinear your data is both qualitatively and quantitatively? How might specific properties of your data influence outputs of various algorithms?
- When you reproduced your clustering experiments on the datasets projected onto the new spaces created by ICA, PCA, and RP, did you get the same clusters as before? Different clusters? Why or why not? Remember to justify why one output might be more interesting when choosing your demonstrations.
- When you re-ran your neural network algorithms were there any differences in performance? Speed? Consider how you might judge differences in performances and include these notes in your discussion.

It might be difficult to generate the same kinds of graphs for this part of the assignment as you did in previous assignments; however, you should come up with some way to describe the kinds of clusters you produce. If you can achieve this visually all the better. However, a note of caution. Figures should remain legible as we are asking for several demonstrations in many sections. Do not try to squish figures together in specific sections where axis labels become 8pt font or less. We are looking for clear and concise demonstration of knowledge and synthesis of results in your demonstrations. Any paper that solely has figures without formal writing will not be graded. Be methodical with your space.

Note: we need to be able to get to your code and your data. Providing entire libraries isn't necessary when a URL would suffice; however, you should at least provide any files you found necessary to change and enough support and explanation so we can reproduce your results on a standard linux machine.

Discretionary Points: In an effort to encourage exploration and utilize theory and practice from the course material, we will award an additional 5 points of extra credit to this assignment to those who go above and beyond the requirements. This does not mean provide every single plot generated, but rather a different piece of analysis or utilization of a novel method from a recent paper to compliment what is asked in any particular section of the report. The discretionary points are optional as you can do extremely well on the assignment without it.

5 Rescoring Criteria

When your assignment is scored, you will receive feedback explaining your errors and successes in some level of detail. This feedback is for your benefit, both on this assignment and for future assignments. It is considered a part of your learning goal to internalize this feedback.

If you are convinced that your score is in error in light of the feedback, you may request a rescore within a week of the score and feedback being returned to you. A rescore request is only valid if it includes an explanation of where the grader made an error.

It is important to note that because we consider your ability to internalize feedback a learning goal, we also assess it. This ability is considered 10 percent of each assignment. We default to assigning you full credit. If you request a rescore and do not receive at least 5 points as a result of the request, you will lose those 10 points.

6 Plagiarism and Proper Citation

The easiest way to fail this class is to plagiarize. **Using the analysis, code or graphs of others in this class is considered plagiarism.** The assignments are designed to force you to immerse yourself in the empirical and engineering side of ML that one must master to be a viable practitioner and researcher. It is important that you understand why your algorithms work and how they are affected by your choices in data and hyperparameters. The phrase "as long as you participate in this journey of exploring, tuning, and analyzing" is key. We take this very seriously and you should too.

What is plagiarism?

If you copy any amount of text from other students, websites, or any other source without proper attribution, that is plagiarism. The most common form of plagiarism is copying definitions or explanations from wikipedia or similar websites. We use an anti-cheat tool to find out which parts of the assignments are your own and there is a near 100 percent chance we will find out if you copy or paraphrase text or plots from online articles, assignments of other students (even across sections and previous courses), or website repositories.

What does it mean to be original?

In this course, we care very much about your analysis. It must be original. Original here means two things: 1) the text of the written report must be your own and 2) the exploration that leads to your analysis must be your own. Plagiarism typically refers to the former explicitly, but in this case it also refers to the latter explicitly.

It is well known that for this course we do not care about code. We are not interested in your working out the edge cases in k-nn, or proving your skills with python. While there is some value in implementing algorithms yourselves in general, here we are interested in your grokking the practice of ML itself. That practice is about the interaction of algorithms with data. As such, the vast majority of what you're going to learn in order to master the empirical practice of ML flows from doing your own analysis of the data, hyper parameters, and so on; hence, you are allowed to steal ML code from libraries but are not allowed to steal code written explicitly for this course, particularly those parts of code that automate exploration. You will be tempted to just run said code that has already been overfit to the specific datasets used by that code and will therefore learn very little.

How to cite:

If you are referring to information you got from a third-party source or paraphrasing another author, you need to cite them right where you do so and provide a reference at the end of the document [Col]. Furthermore, "if you use an author's specific word or words, you must place those words within quotation marks and you must credit the source." [Wis]. It is good style to use quotations sparingly. Obviously, you cannot quote other people's assignment and assume that is acceptable. Speaking of acceptable, citing is not a get-out-of-jail-free card. You cannot copy text willy nilly, but cite it all and then claim it's not plagiarism just because you cited it. Too many quotes of more than, say, two sentences will be considered plagiarism and a terminal lack of academic originality.

Your README file will include pointers to any code and libraries you used.

If we catch you...

We report all suspected cases of plagiarism to the Office of Student Integrity. Students who are under investigation are not allowed to drop from the course in question, and the consequences can be severe, ranging from a lowered grade to expulsion from the program.

7 Version Control

- 01/22/2024 - TJL copied assignment to JM's L^AT_EX template.
- 03/01/2024 - TJL updated class percent total, phrasing, discretionary points.

References

- [Col] Williams College. *Citing Your Sources: Citing Basics*. URL: <https://libguides.williams.edu/citing>.
- [Wis] University of Wisconsin - Madison. *Quoting and Paraphrasing*. URL: <https://writing.wisc.edu/handbook/assignments/quotingsources>.

Original assignment by Charles Isbell. Updated for Spring 2024 by John Mansfield and Theodore LaGrow.
Modified for L^AT_EX by John Mansfield.